AZURE AI / EXAM PRACTICE

AI-103
Practice Set B

Practice with answers

Developing AI Apps and Agents on Azure
100 original questions · 8 October 2026

Coverage	Questions
Plan and manage an Azure AI solution	28
Generative AI and agentic solutions	33
Computer vision	13
Text analysis and speech	13
Information extraction	13

How to use this booklet

Try each question before checking the highlighted solution and explanation. For a separate attempt, use the companion questions-only booklet.

Choose one answer unless the question says to select two. For sequences, arrange the lettered steps. For matching, use the provided choice bank; one choice is unused.

64 single-answer items, including case studies and exhibits; 18 multiple-select; 10 sequences; 8 matching. Case-scenario links connect the relevant questions.

Objectives effective 16 April 2026. Difficulty labels are author estimates. This is independent practice material, not an official exam or a prediction of your exam score.

Q1

Domain A: Plan and manage an Azure AI solution | Single choice | medium

A regulated assistant uses Data Zone Standard in the EU. A proposed switch to Global Standard keeps the resource in the same EU region. What must the architecture review flag?

A Keeping the resource in an EU region also keeps the Global Standard deployment's inference processing inside the same EU boundary.

B Requests must move to an asynchronous batch submission interface.

C Inference processing can extend beyond the EU despite keeping the resource in its existing EU region. ✓

D The model begins billing against a reserved throughput commitment.

Why the correct answer is right

The deployment type determines the processing boundary. Keeping the Azure resource in an EU region does not give a Global deployment the EU Data Zone processing restriction.

Why the other options are wrong

- **A:** The processing boundary follows deployment type, not just resource location.
- **B:** Global Standard is an interactive standard deployment.
- **D:** Global Standard uses standard token billing, not provisioned throughput.

Sources

Deployment types for Microsoft Foundry Models
https://learn.microsoft.com/azure/foundry/foundry-models/concepts/deployment-types

Q2

Domain A: Plan and manage an Azure AI solution | Single choice | easy

A nightly classification job can wait for asynchronous completion. It has no interactive users and wants the batch charging model. Which deployment family should the team assess first?

A Provisioned deployment with continuously reserved inference capacity.

B Developer deployment used as the long-running production job endpoint.

C Standard deployment with synchronous inference.

D Batch deployment for asynchronous inference jobs. ✓

Why the correct answer is right

Batch is intended for asynchronous groups of inference requests. Select its geographic variant and supported model according to the job requirements.

Why the other options are wrong

- **A:** The stated billing and completion requirements fit Batch.
- **C:** This uses the standard interactive path rather than the specified batch completion and charging model.
- **B:** Developer deployment is intended for experimentation.

Sources

Deployment types for Microsoft Foundry Models
https://learn.microsoft.com/azure/foundry/foundry-models/concepts/deployment-types

Q3

Domain A: Plan and manage an Azure AI solution | Single choice | medium

A team migrated a tool service to Entra authentication. Its identity is assigned a narrow role on the wrong Foundry resource, and requests are denied. Which correction addresses the authorization scope?

A Copy the deployment operator role onto a developer account used for manual checks.

B Change the application's trusted origins.

C Grant the same role to the identity on the build pipeline resource group.

D Assign the required role on the resource the service actually calls. ✓

Why the correct answer is right

RBAC is evaluated for the calling principal at the target scope. The role and its scope must cover the resource being accessed.

Why the other options are wrong

- **C:** Permissions on an unrelated resource do not authorize the target.
- **A:** A developer account is not the calling service identity.
- **B:** Trusted origins do not grant RBAC permissions on the target resource.

Sources

Role-based access control for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry

Q4

Domain A: Plan and manage an Azure AI solution | Single choice | medium

An agent resolves Search to its private address and establishes a network connection, but the runtime receives an authorization denial. A developer can query the index. Which evidence is most useful next?

A The developer's successful query together with the developer's subscription permissions.

B The private DNS zone records already confirmed by the agent's connection trace.

C The runtime connection principal and its role assignments on Search. ✓

D The model's token allocation and the embedding model's vector dimensions.

Why the correct answer is right

Network reachability and authorization are separate checks. Inspect the identity actually used by the agent connection rather than inferring access from a developer session.

Why the other options are wrong

- **A:** That proves developer access, not the agent principal permissions.
- **B:** Resolution is established; the observed failure concerns authorization.
- **D:** Those settings do not explain a Search authorization denial.

Sources

Set up private networking for Foundry Agent Service
https://learn.microsoft.com/azure/foundry/agents/how-to/virtual-networks

Role-based access control for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry

Q5

Domain A: Plan and manage an Azure AI solution | Single choice | easy

Managed-compute quota was approved for accelerator SKU X in region R1. A deployment of SKU X in region R2 still fails for capacity quota. What is the relevant check?

A Check the managed-compute accelerator-SKU quota in R2; the R1 approval is scoped to its own region. ✓

B Check the Azure VM quota for the matching GPU family in R2 and count that unused allocation toward the managed-compute accelerator requirement.

C Whether the deployment's API key was issued by a resource in R2.

D Whether the serverless endpoint in R1 has unused model token throughput.

Why the correct answer is right

The allocation is scoped to both the accelerator SKU and region. An approval in R1 does not establish the corresponding managed-compute allocation in R2.

Why the other options are wrong

- **B:** Azure VM quota is separate from the managed-compute allocation.
- **D:** Serverless model throughput is a different limit.
- **C:** The authentication key does not allocate accelerator capacity.

Sources

Deployment overview for Microsoft Foundry Models
https://learn.microsoft.com/azure/foundry/concepts/deployments-overview

Q6

Domain A: Plan and manage an Azure AI solution | Single choice | medium

A hosted agent repeatedly returns regional_session_quota_exceeded despite appropriate backoff. Moving regions is prohibited, and model token capacity is available. Which escalation fits?

A Increase the underlying model token allocation in the current region.

B Open an Azure support request for the persistent regional limit. ✓

C Retry the same request with a shorter delay.

D Delete retained conversation records to reduce the size of the Search index.

Why the correct answer is right

After suitable retry and backoff, a persistent regional session limit calls for support when a regional alternative is not allowed.

Why the other options are wrong

- **A:** The error identifies regional session capacity, not model tokens.
- **C:** Shorter delays do not resolve the persistent regional capacity limit.
- **D:** Search storage and retained records do not diagnose the regional session limit.

Sources

Foundry Agent Service limits, quotas, and regional support
https://learn.microsoft.com/azure/foundry/agents/concepts/limits-quotas-regions

Q7

Domain A: Plan and manage an Azure AI solution | Single choice | easy

A deployment workflow can obtain an Entra token through OIDC, but deploying the hosted agent is denied. Project endpoint variables are correct. What should be checked first?

A The environment variable names used to display the workflow's build number.

B The runtime agent's Search query filters for the target index.

C The prompt token limit of the model called by the deployed agent.

D Deployment permissions of the federated workflow identity. ✓

Why the correct answer is right

Successful federation proves token acquisition, not permission to deploy. Check role assignments on the resource scope used by the workflow.

Why the other options are wrong

- **B:** These control retrieval behavior, not deployment authorization.
- **C:** Model token limits do not authorize deployment operations.
- **A:** Display metadata does not grant the deployment role.

Sources

Quickstart: Hosted agent CI/CD templates
https://learn.microsoft.com/azure/foundry/agents/quickstarts/set-up-cicd-hosted-agent

Q8

Domain A: Plan and manage an Azure AI solution | Single choice | medium

A dashboard shows a low evaluation score, but engineers cannot locate the request trace that produced it. Trace spans already include operation timings. Which addition closes the gap?

A Increase the number of unrelated aggregate latency samples collected each hour.

B Correlate evaluation run identifiers with request trace identifiers. ✓

C Record the evaluator's display name as the identifier for all evaluated requests.

D Rename each tool operation to include its service's public hostname.

Why the correct answer is right

Correlation identifiers let engineers move from an evaluation result to the associated request and operation spans.

Why the other options are wrong

- **A:** Aggregates do not link a scored result to its request.
- **C:** A shared name does not identify a particular evaluation run.
- **D:** Operation names do not establish evaluation-to-trace linkage.

Sources

Agent tracing overview
https://learn.microsoft.com/azure/foundry/observability/concepts/trace-agent-concept

Q9

Domain A: Plan and manage an Azure AI solution | Single choice | medium

A team redacts span content but exports traces to an unrestricted analytics workspace with indefinite retention. Which remaining control is most relevant?

A Increase sampling so the unrestricted telemetry workspace contains a representative trace for each request category and evaluation outcome.

B Apply operational log access restrictions and retention periods to the redacted trace storage. ✓

C Replace the trace identifiers with the model deployment name for easier grouping.

D Move the redaction stage from the exporter into a scheduled cleanup task.

Why the correct answer is right

Redaction is one part of telemetry protection. Trace storage should also follow appropriate access and retention controls.

Why the other options are wrong

- **A:** Sampling does not restrict access or apply the required retention policy.
- **C:** Grouping does not enforce the required controls.
- **D:** That permits sensitive data to be stored before redaction.

Sources

Agent tracing overview
https://learn.microsoft.com/azure/foundry/observability/concepts/trace-agent-concept

Q10 Domain A: Plan and manage an Azure AI solution | Single choice | medium

Search requests intermittently slow down and some are throttled during traffic bursts. Which alert pair directly targets these symptoms?

A Model answer-quality alerts.

B Search latency and throttled-request alerts. ✓

C Speech recognition accuracy and audio-duration alerts.

D Indexer field-confidence and document-page-count alerts.

Why the correct answer is right

Search operational monitoring includes query latency and throttling. Use resource logs and workload context to investigate alerts.

Why the other options are wrong

- **D:** Those do not measure query-serving latency or throttling.
- **A:** Answer quality does not directly measure Search query-serving latency or throttling.
- **C:** These measure a different service and modality.

Sources

Monitor Azure AI Search
https://learn.microsoft.com/azure/search/monitor-azure-cognitive-search

Q11 Domain A: Plan and manage an Azure AI solution | Single choice | easy

A caller supplies a receipt for one support interaction. The agent already has persistent user preferences and a maintained policy knowledge base. Which capability best supports lookup in this receipt?

A Persistent memory extraction used as the receipt's full document retrieval layer.

B File search over the user-provided document. ✓

C The policy knowledge base populated with permanent copies of each temporary receipt.

D The model router configured to select a larger context window.

Why the correct answer is right

File search is intended for user-provided documents during an interaction. It complements rather than replaces memory and curated knowledge.

Why the other options are wrong

- **A:** Memory stores distilled context, not the interaction's document search surface.
- **C:** That treats a temporary upload as curated organizational content.
- **D:** Routing does not itself provide retrieval over the uploaded file.

Sources

Memory in Microsoft Foundry Agent Service (preview)
https://learn.microsoft.com/azure/foundry/agents/concepts/what-is-memory

Q12 Domain A: Plan and manage an Azure AI solution | Single choice | easy

A migration mixes azure-ai-projects 1.0.0 examples with a Foundry (new) project endpoint and 2.x code snippets. What is the most useful first correction?

A Align the package version, endpoint and examples to the 2.x client. ✓

B Change the Search index vector dimensions to match the project identifier.

C Keep the mixed examples and increase the endpoint request timeout.

D Replace the Entra credential with a resource key for the mixed client code.

Why the correct answer is right

Foundry (new) and classic examples use different client generations. A coherent package, endpoint and operation surface is necessary before debugging individual calls.

Why the other options are wrong

- **C:** Timeouts do not reconcile incompatible client generations.
- **B:** Index dimensions are unrelated to SDK compatibility.
- **D:** The specified 2.x client requires Entra authentication.

Sources

Microsoft Foundry SDKs and endpoints
https://learn.microsoft.com/azure/foundry/how-to/develop/sdk-overview

Azure AI Projects client library for Python - version 2.8.0
https://learn.microsoft.com/python/api/overview/azure/ai-projects-readme?view=azure-python

Q13 Domain A: Plan and manage an Azure AI solution | Single choice | easy

An agent obtains the requested result but violates the instruction to use a read-only workflow. Which evaluator most directly addresses that policy failure?

A Task Adherence (preview), assessing compliance with the instruction to use a read-only workflow. ✓

B Tool Call Success, assessing whether an invoked tool completes successfully.

C Tool Output Utilization, assessing use of returned tool information.

D Task Completion (preview), treating successful delivery of the requested result as evidence that the read-only workflow instruction was obeyed.

Why the correct answer is right

Task Adherence evaluates compliance with instructions and constraints. A completed task can still have an adherence failure.

Why the other options are wrong

- **D:** Completing the task does not establish policy adherence.
- **C:** Using results well does not establish policy adherence.
- **B:** A technically successful call can violate the operating policy.

Sources

Agent evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators

Q14 Domain A: Plan and manage an Azure AI solution | Single choice | easy

Two RAG systems retrieve the same passages. One produces a useful answer but adds a claim absent from those passages. Which evaluation best isolates that defect?

A Response Completeness (preview) against the reference answer coverage.

B Relevance against the wording of the user question.

C Groundedness against the supplied passages, which tests the unsupported claim in the generated answer. ✓

D Document Retrieval, using the identical retrieved passages to conclude that any later difference in generated claims must already be context-supported.

Why the correct answer is right

Groundedness measures support for the generated response in its supplied context. It complements retrieval quality and relevance.

Why the other options are wrong

- **D:** Retrieval evaluation does not establish support for the generated claim.
- **B:** An unsupported claim can still address the topic.
- **A:** Coverage does not by itself demonstrate contextual support.

Sources

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Q15 Domain A: Plan and manage an Azure AI solution | Single choice | medium

An agent can call a purchasing tool. The team wants a guardrail to inspect the proposed action before execution. Which intervention point best fits the preview agent controls?

A The tool-call intervention point, configured to inspect the proposed purchasing action before execution. ✓

B The tool-response intervention that scans returned purchase receipts.

C The final-output intervention, configured to inspect the purchase receipt and ask for human approval before displaying the already completed action.

D The user-input scan performed before the model proposes an action.

Why the correct answer is right

A tool-call intervention can inspect a proposed call before execution. Input, tool-response and final-output interventions protect different boundaries.

Why the other options are wrong

- **C:** The purchase has executed before this check; the required inspection must precede execution.
- **B:** That boundary is after execution.
- **D:** It cannot inspect the final proposed tool call and arguments.

Sources

Guardrails and controls overview in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/guardrails/guardrails-overview

Q16 Domain A: Plan and manage an Azure AI solution | Single choice | easy

A security design requires the memory store to integrate with a virtual network. The team proposes the documented preview memory service. What must the design record?

A That using Entra authentication enables the memory store private network interface.

B That the documented preview lacks memory-store VNet integration. ✓

C That placing the caller in a subnet gives the memory store the same network boundary.

D That assigning the project owner role enables preview private network features.

Why the correct answer is right

The documented preview does not support memory-store VNet integration. This is a capability constraint to resolve before choosing that service for the requirement.

Why the other options are wrong

- **C:** Caller placement does not add an unsupported service feature.
- **A:** Authentication does not create VNet integration.
- **D:** Roles do not change the documented capability boundary.

Sources

Memory in Microsoft Foundry Agent Service (preview)
https://learn.microsoft.com/azure/foundry/agents/concepts/what-is-memory

Q17 Domain A: Plan and manage an Azure AI solution | Single choice | easy

A team changes its model deployment to Provisioned but leaves the agent on basic setup. It expects agent state to move into its subscription. Which assessment is correct?

A Adding a private DNS record migrates the existing state into customer-managed storage.

B A Search role assignment on the model endpoint changes the state storage provider.

C Provisioned model throughput also provisions the agent state resources in the subscription.

D Agent setup must change to standard for customer-managed state resources. ✓

Why the correct answer is right

Model deployment capacity and agent setup are distinct. Select standard setup when the state storage requirement calls for customer-managed resources.

Why the other options are wrong

- **C:** These are separate architecture choices.
- **A:** DNS does not migrate state ownership.
- **B:** Authorization on another resource does not select agent state storage.

Sources

Foundry Agent Service limits, quotas, and regional support
https://learn.microsoft.com/azure/foundry/agents/concepts/limits-quotas-regions

Set up private networking for Foundry Agent Service
https://learn.microsoft.com/azure/foundry/agents/how-to/virtual-networks

Q18 Domain A: Plan and manage an Azure AI solution | Single choice | easy

A routing evaluation looks acceptable, but governance approves just two underlying models. The router currently has a broader eligible pool. What should change before deployment?

A Increase the routing evaluation sample until both approved models appear frequently.

B Restrict the routed model subset to the approved models. ✓

C Keep the pool and add the approved model names to the user-facing system prompt.

D Filter final answers whose text mentions an unapproved model name.

Why the correct answer is right

Approval of specific models is enforced by restricting the router subset. Evaluation frequency alone is not a routing constraint.

Why the other options are wrong

- **C:** A prompt is not the supported pool restriction.
- **D:** An answer need not name the model that generated it.
- **A:** Sampling does not exclude unapproved models from future routing.

Sources

Use model router for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/openai/how-to/model-router

Case study 1

Questions 19, 20, 21, 22.

Case study - Atlas service assistant

Answer each question using the shared scenario. Your answers do not change the scenario.

Service context

Atlas sells replacement equipment and operates an internal assistant. Inventory lookups, reservations and shipment or account changes are implemented by host application functions.

Operating rules

Business systems are authoritative for inventory and action outcomes. Tool inputs must be validated against business rules and the caller's account access. Actions requiring approval must be approved before execution. Each question describes a separate observation or proposed implementation.

Q19 Domain B: Implement generative AI and agentic solutions Case study medium

> Use case study 1: Atlas service assistant.

An agent proposes reserve_part, and the host application returns 'success' before contacting the inventory service. A later check finds no reservation. What is the design defect?

A The host should first send the success result back to the model, then execute the reservation if the model confirms that the JSON was valid.

B The function declaration needed to include the reservation result as a default argument.

C The host reported a tool result before executing the requested operation. ✓

D A missing result default in the tool schema.

Why the correct answer is right

A function request is not proof of execution. The host should execute the validated operation and return its actual result.

Why the other options are wrong

- **A:** This still reports success before execution; model confirmation of JSON does not create the reservation.
- **B:** A default argument is not proof that the business operation ran.
- **D:** A schema default does not execute the reservation or establish its outcome.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q20 Domain B: Implement generative AI and agentic solutions Case study medium

> Use case study 1: Atlas service assistant.

A Responses application executes three tool calls concurrently. Results finish in a different order from the requests. Which rule makes result association reliable?

A Identify outputs using the common tool function name.

B Assign an output identifier based on a sorted combination of the function name and arguments, and retain call_id in application logs.

C Assign outputs to requests by completion order.

D Preserve each originating call_id when constructing its function_call_output. ✓

Why the correct answer is right

Correlation follows call_id rather than arrival order. Concurrent completion is safe when each output retains the identifier of the call it answers.

Why the other options are wrong

- **C:** Execution duration does not identify the originating request.
- **A:** Several calls can invoke the same function.
- **B:** The returned function_call_output must itself use the originating call_id.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q21 Domain B: Implement generative AI and agentic solutions Case study medium

> Use case study 1: Atlas service assistant.

A schema validates a tool's JSON arguments, but the operation would modify an account the caller cannot access. Which control is missing?

A Expand the JSON Schema with a description naming the permitted accounts and accept any generated arguments that satisfy its declared types.

B Application authorization for the requested account and action. ✓

C A tool output stating that the requested account exists.

D A schema example with the desired account and a successful result.

Why the correct answer is right

A schema can constrain the argument structure; it does not decide whether this caller may perform the action on the target account.

Why the other options are wrong

- **A:** A descriptive schema does not enforce authorization for the actual caller and target account.
- **D:** An example does not authorize the current caller.
- **C:** Existence does not establish permission for the caller to modify it.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q22 Domain B: Implement generative AI and agentic solutions Case study hard

> Use case study 1: Atlas service assistant.

A user approves shipping two replacement units. Before execution, the proposed tool arguments change to twenty units. What should the host do?

A Use the earlier approval token with the new argument payload.

B Execute and include the changed quantity in the completion notice.

C Execute because the function name still matches the approved tool.

D Validate the changed proposal and obtain approval for the revised quantity. ✓

Why the correct answer is right

Bind confirmation to the action being executed. A material argument change creates a revised proposal that needs validation and approval.

Why the other options are wrong

- **C:** Approval covered the specific proposal, not any use of the function.
- **B:** Post-execution notice does not approve the revised side effect.
- **A:** The earlier approval does not cover changed business parameters.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Case study 2

Questions 23, 24, 25, 26.

Case study - Meridian knowledge assistant

Answer each question using the shared scenario. Your answers do not change the scenario.

Knowledge workflow

Meridian maintains equipment manuals and policy documents in Azure AI Search. Its assistant retrieves evidence, drafts an answer and may call a business lookup tool.

Review scope

Engineers assess retrieval, generation and tool behavior separately. The following questions describe independent observations and proposed changes; one answer does not change the configuration stated in another question.

Q23 Domain B: Implement generative AI and agentic solutions Case study medium

Use case study 2: Meridian knowledge assistant.

An engineer says enabling semantic reranking creates embeddings and replaces the vectorization stage. Which correction is accurate?

A The semantic reranker creates a numerical representation from the merged scores that can replace the stored document vectors on later queries.

B Reranking adjusts retrieved rankings; an embedding model still creates vectors. ✓

C Reranking converts keyword scores into index-time embeddings.

D Reranking requires the client to sort results by document identifier.

Why the correct answer is right

Vectorization, initial retrieval and semantic reranking are different stages. A semantic reranker evaluates candidate results rather than replacing the embedding model.

Why the other options are wrong

- **A:** Reranker scores are not index-time embeddings and do not replace the content vectorization stage.
- **C:** It does not produce the indexed vector representation.
- **D:** Client identifier sorting is not semantic reranking.

Sources

Hybrid search using vectors and full-text search in Azure AI Search
https://learn.microsoft.com/azure/search/hybrid-search-overview

Q24

Domain B: Implement generative AI and agentic solutions | Case study | medium

> Use case study 2: Meridian knowledge assistant.

A team replaces its embedding model with one using a different representation. It updates the query vectorizer but leaves existing index vectors unchanged. Which migration step is needed?

A Rename the vector field to match the new model deployment name.

B Retain the old vectors and increase the query retry limit.

C Re-embed the indexed content with the compatible model and align index configuration. ✓

D Keep the stored vectors and add the old model name to the query metadata so the service can reinterpret the new model vectors in the old space.

Why the correct answer is right

A model change must keep indexed and query embeddings compatible. Recompute the content vectors and update field dimensions when the new model requires it.

Why the other options are wrong

- **D:** Metadata does not transform representations between unrelated embedding models.
- **B:** Retries preserve the representation mismatch.
- **A:** Field names do not change numerical embeddings.

Sources

Integrated vector embedding in Azure AI Search
https://learn.microsoft.com/azure/search/vector-search-integrated-vectorization

Q25

Domain B: Implement generative AI and agentic solutions | Case study | medium

Use case study 2: Meridian knowledge assistant.

A RAG answer accurately quotes the supplied warranty text but answers a question about delivery lead times. What should be evaluated to capture the mismatch?

A Response relevance to the user's delivery question. ✓

B The count of citations included in the final answer.

C Groundedness as proof of answering the requested topic.

D Embedding-model quota as the measure of answer usefulness.

Why the correct answer is right

Relevance measures how the answer addresses the query. A grounded answer can still respond to the wrong topic.

Why the other options are wrong

- **C:** Support in context does not establish responsiveness to the question.
- **D:** Quota is an operational constraint, not this response-quality measure.
- **B:** Citation count does not show that the requested topic was answered.

Sources

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Q26

Domain B: Implement generative AI and agentic solutions | Case study | medium

> Use case study 2: Meridian knowledge assistant.

The stock lookup tool returns 12 units, but the agent's answer reports 30. The chosen tool and its input arguments were correct. Which evaluator most directly targets the failure?

A Tool Selection for choosing the inventory lookup tool.

B Tool Output Utilization for use of the returned stock information. ✓

C Tool Call Success, using the successful inventory response as evidence that the final answer faithfully represents its returned quantity.

D Tool Input Accuracy for the arguments passed to the tool.

Why the correct answer is right

Tool Output Utilization evaluates how the agent uses returned information. A successful and correctly formed call does not ensure correct use of its result.

Why the other options are wrong

- **A:** Tool selection was correct.
- **D:** The arguments were correct.
- **C:** Successful execution does not establish how the agent uses or reports the result.

Sources

Agent evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators

Q27 Domain B: Implement generative AI and agentic solutions Single choice easy

AIProjectClient is authenticated and connected to a Foundry (new) project. The application next needs that project's OpenAI client for Responses operations. Which method fits?

A Call responses.create() directly on the project connection manager.

B Use the project's storage client as the Responses transport.

C Construct an OpenAI client using the Search connection endpoint.

D Call get_openai_client() on the project client. ✓

Why the correct answer is right

The project client provides get_openai_client() for supported OpenAI operations, including the Responses workflow.

Why the other options are wrong

- **A:** Project connection management is not the OpenAI Responses client surface.
- **C:** The Search endpoint is not the project-provided inference client endpoint.
- **B:** A storage client does not implement the Responses API.

Sources

Azure AI Projects client library for Python - version 2.8.0
https://learn.microsoft.com/python/api/overview/azure/ai-projects-readme?view=azure-python

Q28 Domain B: Implement generative AI and agentic solutions Configuration easy

A service migrates a strict JSON-output request from Responses to Chat Completions. It retains text.format and gets a request-shape error. Which field should carry the format setting in the destination API?

Migrated request fragment

```
POST .../chat/completions
{"text":{"format":{"type":"json_schema", "...":"..."}}}
```

A response_format, with the supported JSON Schema configuration. ✓

B metadata.format, alongside application tracking labels.

C input.format, wrapping the original input message.

D Keep the format configuration inside text.format and select the Chat Completions endpoint, relying on the identical schema to make the request portable.

Why the correct answer is right

The destination API has a different format location: response_format for Chat Completions. Retain a supported schema and adapt the request contract.

Why the other options are wrong

- **D:** The schema can be reused when supported, but the request location must follow the destination API contract.
- **C:** This is not the documented structured-output field.
- **B:** Metadata labels do not enforce the response format.

Sources

Structured outputs
https://learn.microsoft.com/azure/foundry/openai/how-to/structured-outputs

Q29 Domain B: Implement generative AI and agentic solutions Single choice easy

A model must return exactly the defined object fields under the documented strict structured-output subset. Which schema setting rejects additional properties?

A Set additionalProperties to false on the object schema. ✓

B Set additionalProperties to true.

C Rename the schema title to the list of permitted property names.

D Set required to the known names and leave extra properties unrestricted.

Why the correct answer is right

additionalProperties: false restricts the object to declared properties. Required fields and nullable values are configured separately.

Why the other options are wrong

- **B:** That explicitly permits additional properties.
- **D:** Required controls presence, not permission for extra fields.
- **C:** A title does not constrain the object's additional fields.

Sources

Structured outputs
https://learn.microsoft.com/azure/foundry/openai/how-to/structured-outputs

Q30 Domain B: Implement generative AI and agentic solutions Single choice easy

A strict structured-output client assumes every response contains a business object. A safety refusal arrives, and its parser crashes. Which handling is needed?

A Treat the refusal string as the value of the first required business field.

B Retry with a longer property description and parse the same response again.

C Remove additionalProperties: false from the business schema.

D Detect refusal outcomes before parsing a successful structured result. ✓

Why the correct answer is right

Strict output guarantees apply to successful structured responses within the API contract. A client must distinguish refusals from those results.

Why the other options are wrong

- **A:** A refusal is not the requested business output.
- **C:** Relaxing the schema does not turn a refusal into a valid success.
- **B:** Reparsing does not change the refusal outcome.

Sources

Structured outputs
https://learn.microsoft.com/azure/foundry/openai/how-to/structured-outputs

Q31 Domain B: Implement generative AI and agentic solutions Single choice easy

A strict tool schema requires customer_id and permits no extra fields. The client also enables parallel tool calls. Which configuration should be reconciled with the documented structured-output guidance?

A Remove customer_id from required and retain parallel calls.

B Rename customer_id to make the function name unique.

C Change additionalProperties to true for simultaneous arguments.

D Disable parallel_tool_calls while preserving the strict tool schema. ✓

Why the correct answer is right

Use the supported strict-function configuration, including parallel_tool_calls: false. Do not weaken business schema requirements to accommodate an incompatible setting.

Why the other options are wrong

- **A:** That changes the schema requirement rather than the parallel-call setting.
- **C:** This weakens the schema and does not follow the guidance.
- **B:** Property naming does not determine parallel-call support.

Sources

Structured outputs
https://learn.microsoft.com/azure/foundry/openai/how-to/structured-outputs

Q32 Domain B: Implement generative AI and agentic solutions Single choice medium

An engineer wants to diagnose tool latency and is given a model reasoning summary. Which additional evidence directly supports timing analysis?

A A longer reasoning summary described as a full execution trace.

B A list of tool names sorted alphabetically in the prompt.

C Request traces with timed spans for the actual tool operations. ✓

D A paraphrase of the final answer in a second language.

Why the correct answer is right

Use actual operation telemetry for latency diagnosis. A model reasoning summary provides a different kind of information from timed request traces.

Why the other options are wrong

- **A:** A summary is not operation-level telemetry.
- **D:** This does not provide timings.
- **B:** Names alone do not describe runtime duration and relationships.

Sources

Azure OpenAI reasoning models
https://learn.microsoft.com/azure/foundry/openai/how-to/reasoning

Agent tracing overview
https://learn.microsoft.com/azure/foundry/observability/concepts/trace-agent-concept

Q33 Domain B: Implement generative AI and agentic solutions Single choice easy

After switching to a reasoning model, a request rejects a generation parameter used by the previous model. What is the best first step?

A Increase the unsupported parameter until the API accepts its value.

B Set the unsupported control through the old model's deployment metadata.

C Rename the parameter in a prompt while sending the old request field.

D Check the new model's supported parameters and adapt the request contract. ✓

Why the correct answer is right

A model migration must account for supported controls. Check the particular model and API rather than assuming parameter parity.

Why the other options are wrong

- **A:** An unsupported parameter is not made valid by changing its magnitude.
- **C:** The service validates request fields independently of prompt wording.
- **B:** Old deployment metadata does not define the new model request contract.

Sources

Prompt engineering techniques
https://learn.microsoft.com/azure/foundry/openai/concepts/prompt-engineering

Q34 Domain B: Implement generative AI and agentic solutions Single choice medium

A system scans user inputs and tool responses, but the model's final answer may still contain harmful generated content. Which additional intervention covers that output boundary?

A A final-output guardrail before the answer reaches the user. ✓

B A quota check on the model deployment after generation.

C A second scan of the original input message alone.

D A tool-call scan applied to requests with no tool invocation.

Why the correct answer is right

Final-output intervention evaluates the generated answer before delivery. It complements controls on inputs and tools.

Why the other options are wrong

- **C:** This does not inspect the generated answer.
- **D:** It does not cover the final answer boundary.
- **B:** Capacity accounting does not inspect generated content.

Sources

Guardrails and controls overview in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/guardrails/guardrails-overview

Q35 Domain B: Implement generative AI and agentic solutions | Single choice | medium

A Responses-based agent wants to enable Spotlighting using the documented preview. The team assumes a configuration flag will provide feature parity with Chat Completions. What should be checked?

A Whether citations in the final answer demonstrate that Spotlighting was applied.

B The feature's API support boundary before choosing that integration. ✓

C Whether enabling structured outputs activates the Spotlighting feature.

D Whether shortening encoded documents removes the API support restriction.

Why the correct answer is right

The referenced preview supports Spotlighting through Chat Completions. A Responses integration cannot assume the same feature is available.

Why the other options are wrong

- **D:** Smaller inputs do not enable a feature on an unsupported API.
- **C:** Structured output formatting is a separate capability.
- **A:** Citation output does not establish which input protection ran.

Sources

Prompt Shields in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/openai/concepts/content-filter-prompt-shields

Q36 Domain B: Implement generative AI and agentic solutions | Single choice | medium

A Search tool must retrieve documents from a particular department in the configured index. Which capability most directly narrows the returned content?

A Keep retrieval unfiltered and place a department instruction in the system message so the model chooses which returned passages to use.

B A larger vector candidate count for the same unfiltered query.

C A retrieval filter using the department field supported by the index. ✓

D A deployment display name containing the department identifier.

Why the correct answer is right

Retrieval filters constrain the content returned by Search. The index must support the field used in the filter; authorization requirements remain a separate design concern.

Why the other options are wrong

- **A:** This asks the model to ignore content after retrieval rather than narrowing the returned results.
- **B:** More candidates do not enforce the department scope.
- **D:** Display metadata does not scope Search results.

Sources

Connect an Azure AI Search index to Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/ai-search

Q37 Domain B: Implement generative AI and agentic solutions Single choice easy

The correct inventory tool is selected with correct arguments, but it returns a service error and no result. Which signal best describes the immediate failure?

A Tool Input Accuracy, because any service error proves invalid arguments.

B Tool Output Utilization, because the agent misquoted a returned count.

C Tool Selection, because any failed call implies an unsuitable tool.

D Tool Call Success, because the invocation did not complete successfully. ✓

Why the correct answer is right

Invocation success is distinct from selection, input correctness and use of returned data. The immediate defect is the unsuccessful tool call.

Why the other options are wrong

- **C:** The selected capability was appropriate.
- **A:** The scenario establishes correct arguments.
- **B:** No count was returned to misquote.

Sources

Agent evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators

Q38 Domain B: Implement generative AI and agentic solutions Single choice easy

A developer expects a new conversation to retain a preference stored only in the previous session's conversation state. Which capability is needed for persistent interaction-derived context?

A A Search index containing the application's static welcome message.

B A new session identifier with the same human-readable title.

C A configured memory service that retains distilled context across sessions. ✓

D A longer response token budget on the new conversation.

Why the correct answer is right

Session state provides short-term conversational context. Persistent memory is the distinct capability for learned information retained across sessions.

Why the other options are wrong

- **D:** Token budget does not restore previous session context.
- **A:** That does not store the learned user preference.
- **B:** A title does not carry prior state into a new session.

Sources

Memory in Microsoft Foundry Agent Service (preview)
https://learn.microsoft.com/azure/foundry/agents/concepts/what-is-memory

Q39 Domain B: Implement generative AI and agentic solutions Single choice easy

A team narrows the router's model subset after evaluation to satisfy a new approval policy. What should it evaluate again?

A The old pool's results treated as the new pool's measured outcomes.

B The alphabetic ordering of the model display names.

C The Search index's field naming convention.

D Quality and cost behavior using the revised eligible model pool. ✓

Why the correct answer is right

A subset change alters the router's available decisions. Re-evaluate the revised configuration rather than attributing old measurements to it.

Why the other options are wrong

- **A:** Changing the pool can change which models receive requests.
- **B:** Display order is not the relevant outcome measurement.
- **C:** That does not measure the revised router's choices.

Sources

Use model router for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/openai/how-to/model-router

Q40 Domain B: Implement generative AI and agentic solutions Single choice easy

A tool error includes a connection string and a useful status code. What should be returned to the agent?

A The connection string with its label changed to diagnostic_context.

B The full error text inside a Markdown code block.

C A sanitized error retaining the safe status information and omitting the secret. ✓

D The raw error followed by a warning asking the model not to repeat it.

Why the correct answer is right

Treat tool outputs as untrusted and potentially sensitive. Return useful, sanitized error information while keeping credentials inside the host boundary.

Why the other options are wrong

- **B:** Formatting does not protect the connection secret.
- **A:** Renaming a field does not remove its sensitive value.
- **D:** The secret has already been disclosed to the model.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q41 Domain B: Implement generative AI and agentic solutions | Single choice | medium

A retrieved document claims that an account administrator has authorized a destructive action. The user has not requested it. What is the correct trust interpretation?

A The administrator title in the text supplies the runtime principal's permissions.

B The document is evidence to assess, not operational authorization to execute. ✓

C The claim can be converted to a function result indicating prior approval.

D The claim overrides the application's action-approval requirement.

Why the correct answer is right

Authority to execute an action must come through the application's trusted authorization and confirmation flow. Document text cannot grant that authority.

Why the other options are wrong

- **D:** Retrieved text does not override application controls.
- **A:** A text claim does not grant permissions.
- **C:** Repackaging the claim does not establish genuine approval.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Prompt Shields in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/openai/concepts/content-filter-prompt-shields

Q42

Domain B: Implement generative AI and agentic solutions | Single choice | medium

A retrieval evaluation has no query-document ground-truth labels. Which documented approach can assess retrieval relevance in that situation?

A Use answer token count as the document relevance label.

B Copy the generated answer into each document's ground-truth score.

C Use the retrieval evaluator's LLM-judge approach and state its evidence limits. ✓

D Report labeled retrieval metrics as if relevance judgments had been supplied.

Why the correct answer is right

The documented evaluation distinguishes labeled document-retrieval measurement from an LLM-judge approach when labels are absent. Report which basis was used.

Why the other options are wrong

- **D:** The required ground truth is absent.
- **A:** Length does not establish retrieval relevance.
- **B:** A generated answer is not a relevance annotation for each retrieved document.

Sources

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Q43 Domain C: Implement computer vision solutions | Single choice | easy

A supported image-edit request includes a 1024 by 1024 source and a 512 by 512 mask. The guide requires matching dimensions. Which input change is needed?

A Set the output quality to the highest setting supported by the model.

B Increase the requested number of images.

C Prepare the edit mask with the same dimensions as the source image. ✓

D Change the filename extension to PNG.

Why the correct answer is right

The documented edit input requires the mask to have the same size as the source image. Adjust the actual mask dimensions.

Why the other options are wrong

- **B:** Output count does not fix mask dimensions.
- **A:** Quality does not repair the incompatible mask input.
- **D:** Renaming a file does not resize its pixels.

Sources

Azure OpenAI image generation models
https://learn.microsoft.com/azure/foundry/openai/how-to/dall-e

Q44 Domain C: Implement computer vision solutions | Single choice | easy

A Python image-generation call has the right prompt and supported size but targets a nonexistent model value. The Azure deployment has a custom name. What belongs in model for the documented call?

A The image endpoint's API version string.

B The base model catalog identifier when its deployment has a different name.

C The Azure image deployment name used by that resource. ✓

D The resource group name containing the image-generation deployment.

Why the correct answer is right

The documented client.images.generate call uses the Azure deployment name. Other request controls still need to be supported by that deployment.

Why the other options are wrong

- **B:** The documented Azure call selects the deployment by its deployed name.
- **D:** A resource group is not the inference deployment name.
- **A:** API version and model deployment are separate request settings.

Sources

Azure OpenAI image generation models
https://learn.microsoft.com/azure/foundry/openai/how-to/dall-e

Q45 Domain C: Implement computer vision solutions | Single choice | easy

A request uses the documented vision chat API with a limit of ten input images. The app wants to submit twelve photographs in one request. Which change respects that input contract?

A Put two image URLs into each of six text fields.

B Reduce the final answer's token limit.

C Keep twelve image entries and concatenate their captions into a shorter text prompt.

D Divide the photographs into requests within the image-count limit. ✓

Why the correct answer is right

Respect the image-count limit of the selected API and model. Splitting the input changes the analysis context, so combine results deliberately if needed.

Why the other options are wrong

- **B:** Output length does not change the number of images accepted.
- **C:** Shorter text does not reduce the count of supplied images.
- **A:** Embedding URLs as plain text is not supplying fewer supported image inputs.

Sources

Use vision-enabled chat models
https://learn.microsoft.com/azure/foundry/openai/how-to/gpt-with-vision

Q46 Domain C: Implement computer vision solutions | Single choice | easy

A decorative wave image adds no information and is not a link or control. How should its HTML alt attribute be set?

A Use alt="" to identify the image as decorative. ✓

B Omit the alt attribute.

C Use alt="wave-header.png".

D Use alt="Decorative blue wave across the entire top of the page".

Why the correct answer is right

An empty alt attribute marks a purely decorative image. This differs from omitting the attribute.

Why the other options are wrong

- **B:** Omission does not explicitly mark the image as decorative.
- **D:** A decorative detail does not need a spoken description.
- **C:** The source filename is not useful alternative content.

Sources

W3C WAI An alt Decision Tree
https://www.w3.org/WAI/tutorials/images/decision-tree/

Q47 Domain C: Implement computer vision solutions | Configuration | medium

Image analysis returns Hate 0, SelfHarm 2, Sexual 0 and Violence 4. The app rejects any Medium or High category. What is the correct decision?

Returned severity values

```
{"Hate":0,"SelfHarm":2,"Sexual":0,"Violence":4}
```

A Reject because SelfHarm 2 is High.

B Allow because the API request returned a successful HTTP status.

C Reject because Violence reaches the policy's Medium threshold. ✓

D Allow because the average score is below 4.

Why the correct answer is right

Violence 4 is sufficient for rejection under the policy. Keep application policy decisions separate from successful API execution.

Why the other options are wrong

- **D:** The policy is based on any category, not an average.
- **B:** Successful analysis does not imply content approval.
- **A:** Severity 2 does not represent High in this contract.

Sources

QuickStart: Analyze image content
https://learn.microsoft.com/azure/ai-services/content-safety/quickstart-image

Q48 Domain C: Implement computer vision solutions | Single choice | medium

A client asks the image harm endpoint to provide a groundedness judgment for a generated textual answer. What is the API selection problem?

A Groundedness detection is a text capability, not image harm classification. ✓

B The client needs a larger image rather than a different capability.

C The response becomes grounded when each image harm category has severity zero.

D The image endpoint's HTTP status supplies the groundedness result.

Why the correct answer is right

Use the capability appropriate to the input and assessment. Image harm categories and textual groundedness answer different questions.

Why the other options are wrong

- **B:** Image dimensions do not enable textual groundedness detection.
- **C:** Visual harm levels do not measure factual support.
- **D:** Request success is not a textual support judgment.

Sources

Azure AI Content Safety Frequently Asked Questions
https://learn.microsoft.com/en-us/azure/ai-services/content-safety/faq

Q49 Domain C: Implement computer vision solutions | Single choice | medium

A video analyzer misses small text on a distant label. The documented preview resizes sampled frames to 512 by 512 pixels. Which explanation should guide follow-up?

A The field schema's title determines the video's original pixel resolution.

B Increasing segment count necessarily restores the lost text pixels.

C A confidence value can reconstruct text absent from the analyzed frame.

D Resizing can remove small details needed to read the label. ✓

Why the correct answer is right

The documented sampling and resizing limits can hide small or distant details. Verify against a suitable original frame or another input strategy rather than inferring unread text.

Why the other options are wrong

- **A:** Schema naming does not control the original visual information.
- **B:** Segmentation does not guarantee preservation of fine detail.
- **C:** Confidence does not recover lost image information.

Sources

Azure Content Understanding in Foundry Tools video solutions
https://learn.microsoft.com/azure/ai-services/content-understanding/video/overview

Q50 Domain C: Implement computer vision solutions | Single choice | easy

A sample uses Content Understanding Python SDK 1.1.0, tied to the 2025-11-01 API. A developer copies a preview-only analyzer feature into it. What should be resolved first?

A Compatibility between the SDK, API version and analyzer feature. ✓

B The ordering of imports in the Python source file.

C The output PDF filename used after analysis.

D The model's natural-language explanation of the preview feature.

Why the correct answer is right

Treat SDK version, API contract and feature availability as one compatibility decision. A GA-bound client should not silently inherit preview-only examples.

Why the other options are wrong

- **B:** Import order does not supply a missing API feature.
- **C:** A filename does not affect service contract support.
- **D:** An explanation does not enable an unsupported request field.

Sources

Azure AI Content Understanding client library for Python - version 1.1.0
https://learn.microsoft.com/python/api/overview/azure/ai-contentunderstanding-readme?view=azure-python

What is Azure Content Understanding in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/content-understanding/overview

Q51 Domain D: Implement text analysis solutions | Single choice | easy

An application must label organizations and locations in a paragraph. It does not need identity redaction or language routing. Which prebuilt capability fits?

A Language detection to select a translated version of the paragraph.

B PII detection focused on personal identifiers.

C Named entity recognition using predefined entity categories. ✓

D Abstractive summarization of the paragraph into a single sentence.

Why the correct answer is right

Prebuilt NER categorizes mentions such as organizations and locations. It serves a different purpose from PII detection, language detection and summarization.

Why the other options are wrong

- **B:** The task is general organization and location labeling.
- **A:** Language routing does not label entity mentions.
- **D:** A summary is not structured entity labeling.

Sources

What is Azure Language in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/language-service/overview

Q52

Domain D: Implement text analysis solutions | Single choice | easy

A product needs a short paraphrased summary that combines related ideas across a document. Which summarization method targets that output?

A Abstractive summarization that generates a compact rewording. ✓

B PII detection followed by concatenation of the found spans.

C Extractive summarization limited to selected source sentences.

D Language detection applied separately to each paragraph.

Why the correct answer is right

Abstractive summarization produces a new wording. Review factual support because a fluent paraphrase is not evidence of fidelity by itself.

Why the other options are wrong

- **C:** This preserves sentences rather than producing the requested paraphrase.
- **D:** Language identification does not summarize the content.
- **B:** Personal-information spans do not form a thematic summary.

Sources

What is summarization?
https://learn.microsoft.com/azure/ai-services/language-service/summarization/overview

Q53 Domain D: Implement text analysis solutions | Single choice | easy

A team needs to translate a collection of documents as a document workflow rather than individual chat messages. Which service surface is the most appropriate starting point?

A Detect each document language and save that label as its translated output.

B Translator document translation, choosing a supported synchronous or asynchronous mode. ✓

C Generate a summary in the target language and return it as the translated document.

D Translate extracted filenames through the text translation endpoint.

Why the correct answer is right

Translator provides document translation modes. Choose the mode based on the supported document types and workload requirements.

Why the other options are wrong

- **A:** Language detection does not translate the document.
- **D:** Translating names does not translate document contents.
- **C:** A summary omits document content and is not a document translation.

Sources

What is Azure Translator in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/translator/overview

Q54 Domain D: Implement text analysis solutions | Single choice | easy

A custom recognition model performs well in evaluation, but the app still calls the base recognizer endpoint. Which step makes the selected custom model available to the application?

A Deploy the model to a custom endpoint and configure the app to use it. ✓

B Set the model ID as the synthesis voice name.

C Upload the evaluation report as the next recognition audio input.

D Rename the training dataset to the application's endpoint hostname.

Why the correct answer is right

Training and evaluation do not switch production traffic. Deploy the chosen recognition model to a custom endpoint and use that endpoint from the app.

Why the other options are wrong

- **B:** Recognition deployment is not a text-to-speech voice selection.
- **D:** Dataset naming does not deploy or route inference.
- **C:** A report does not change the serving model.

Sources

What is custom speech?
https://learn.microsoft.com/azure/ai-services/speech-service/custom-speech-overview

Q55 Domain D: Implement text analysis solutions | Single choice | easy

An Italian-to-German Speech translation call succeeds. The app needs the German text for a later synthesis step. Which result entry should it read?

A The audio output of a SpeechSynthesizer that has not yet been called.

B The translations entry keyed by de, the configured target language. ✓

C The translations entry keyed by it-IT.

D The recognition language field from the request configuration.

Why the correct answer is right

Read the text associated with the configured target language, then pass that text to synthesis if audio output is required.

Why the other options are wrong

- **C:** That is the recognition source, not the German target identifier.
- **A:** Translation success does not mean a separate synthesis step has run.
- **D:** That field describes the input language rather than the translated text.

Sources

How to recognize and translate speech
https://learn.microsoft.com/azure/ai-services/speech-service/how-to-translate-speech

Q56 Domain D: Implement text analysis solutions | Configuration | easy

A prototype works with one Voice Live model in one region. The team plans a new regional deployment with a different model and assumes all voice features transfer. What should be checked?

Deployment change

```
Prototype: model M1, region R1
Proposed:  model M2, region R2
```

A The chosen model, feature and regional compatibility for Voice Live. ✓

B Use a common endpoint hostname as the basis for regional feature compatibility.

C Whether using the previous model's name as an alias enables its old features.

D Whether client-side voice labels can override model compatibility limits.

Why the correct answer is right

A successful prototype establishes behavior for its tested configuration. Voice Live features depend on the selected model and regional support.

Why the other options are wrong

- **C:** An alias does not establish feature support in the new configuration.
- **B:** A common endpoint pattern does not establish regional capability parity.
- **D:** Client labels do not add service capabilities.

Sources

Voice Live API for real-time voice agents
https://learn.microsoft.com/azure/ai-services/speech-service/voice-live

Q57 Domain D: Implement text analysis solutions | Single choice | medium

A team can connect to a realtime audio endpoint but users experience overlapping speech and inconsistent turn endings. What is the most relevant next design review?

A The connection URL after a successful transport handshake.

B The complete-session transcript formatting.

C Session turn detection and how interruption affects input and output audio. ✓

D The static voice chosen for synthesized speech.

Why the correct answer is right

Transport connectivity is one part of realtime audio. Review turn detection, cancellation or interruption, and client audio handling as a coordinated interaction.

Why the other options are wrong

- **D:** Voice identity does not define turn-ending and interruption behavior.
- **B:** Formatting a transcript does not coordinate input and output audio.
- **A:** An established transport does not itself define correct turn management.

Sources

Use the GPT Realtime API for speech and audio
https://learn.microsoft.com/azure/foundry/openai/how-to/realtime-audio

Q58 Domain E: Implement information extraction solutions | Single choice | easy

A source contains a field named legacy_title and the index calls it title. No skill creates or transforms this value. Which mapping is appropriate?

A A semantic reranking query requesting the missing title.

B A vectorizer setting using title as the model deployment name.

C An outputFieldMappings entry pointing at a nonexistent skill output.

D A fieldMappings entry from legacy_title to title. ✓

Why the correct answer is right

Use fieldMappings for source-to-index field mapping. Use outputFieldMappings for values produced in the enrichment tree.

Why the other options are wrong

- **C:** The value is an original source field, not a generated enrichment node.
- **B:** A vectorizer does not rename this source field.
- **A:** A query does not configure ingestion mappings.

Sources

Skillset concepts in Azure AI Search
https://learn.microsoft.com/azure/search/cognitive-search-working-with-skillsets

Q59 Domain E: Implement information extraction solutions | Single choice | medium

OCR receives normalized images and produces text, but the index's ocr_text field remains empty. Which remaining configuration should be inspected?

A The output mapping from the enriched OCR text node to ocr_text. ✓

B The Speech translation target language.

C The maximum response length used by the chat application.

D The imageAction setting already producing normalized images.

Why the correct answer is right

Successful enrichment is distinct from writing its output into the index. Inspect the enriched node path and outputFieldMappings target.

Why the other options are wrong

- **D:** The scenario establishes successful OCR input and output.
- **C:** Generation length does not populate index fields.
- **B:** Speech configuration does not map OCR enrichment output.

Sources

OCR cognitive skill
https://learn.microsoft.com/azure/search/cognitive-search-skill-ocr

Skillset concepts in Azure AI Search
https://learn.microsoft.com/azure/search/cognitive-search-working-with-skillsets

Q60 Domain E: Implement information extraction solutions | Single choice | medium

An integrated vectorization run repeatedly fails when the embedding deployment's quota is exhausted. Retries are configured correctly. What should be addressed?

A The answer generation temperature used after retrieval.

B Embedding capacity or ingestion demand relative to that quota. ✓

C Increase the retry budget and keep the current ingestion concurrency so the queued workload can eventually consume more tokens than the assigned quota.

D The order of fields displayed in the Search portal.

Why the correct answer is right

Integrated vectorization still calls an embedding deployment subject to quota. Persistent exhaustion requires capacity or workload changes, not retries alone.

Why the other options are wrong

- **C:** A retry budget does not increase a persistently exhausted quota; demand or allocated capacity must change.
- **A:** Temperature does not increase embedding throughput.
- **D:** Display order does not affect embedding quota.

Sources

Integrated vector embedding in Azure AI Search
https://learn.microsoft.com/azure/search/vector-search-integrated-vectorization

Q61 Domain E: Implement information extraction solutions | Single choice | easy

A Search enrichment pipeline needs section structure and Markdown from document layout, with image-location metadata available for downstream use. Which skill is designed for this?

A A semantic reranker operating after documents are indexed.

B A text-splitting skill over the raw document text alone.

C A query-time vectorizer configured on the document title.

D The Document Layout skill using Document Intelligence layout analysis. ✓

Why the correct answer is right

The Document Layout skill emits structured Markdown or text from layout analysis and can preserve image-location information.

Why the other options are wrong

- **B:** Splitting does not recover the page layout and image-location metadata.
- **A:** Reranking is a query stage, not layout extraction.
- **C:** Query vectorization does not extract layout from source documents.

Sources

Document Layout skill
https://learn.microsoft.com/azure/search/cognitive-search-skill-document-intelligence-layout

Q62 Domain E: Implement information extraction solutions Single choice easy

An analysis operation is queued. A worker must index the extracted content after successful completion. Which item should it retain while waiting?

A The operation handle needed to poll that submitted analysis. ✓

B The final index document count as a substitute for operation state.

C A fabricated empty content result to mark the document as processed.

D Submit the document as a new analysis on each polling interval and select whichever operation first returns a completed content object.

Why the correct answer is right

Track the submitted long-running operation through completion. The downstream index should receive the actual completed analysis result.

Why the other options are wrong

- **D:** Track the existing operation; repeated submissions create extra analysis work and do not poll the original request.
- **C:** That claims completion before extraction succeeds.
- **B:** Aggregate index count does not track this analysis job.

Sources

Azure AI Content Understanding client library for Python - version 1.1.0
https://learn.microsoft.com/python/api/overview/azure/ai-contentunderstanding-readme?view=azure-python

Q63 Domain E: Implement information extraction solutions | Single choice | medium

A team enables estimateFieldSourceAndConfidence for document fields. It receives a high-confidence amount and wants to describe it as independently verified. Which statement is accurate?

A The score replaces page and bounding-box references for evidence retrieval.

B The score proves that the downstream payment was authorized.

C The score estimates analyzer confidence; source review can still verify the value. ✓

D The score is a guarantee that the amount matches the original document.

Why the correct answer is right

Confidence is a model-derived estimate useful for routing. Grounding allows evidence inspection; neither should be confused with operational approval.

Why the other options are wrong

- **D:** Confidence is not a correctness guarantee.
- **A:** Confidence does not locate the source.
- **B:** Extraction confidence is not action authorization.

Sources

Validate document analyzer quality with confidence, grounding, and labeled samples
https://learn.microsoft.com/azure/ai-services/content-understanding/document/analyzer-improvement

Q64 Domain E: Implement information extraction solutions Single choice easy

A document analyzer performs well on established suppliers but misreads amounts on a new invoice layout. What input can help improve it after schema definitions are reviewed?

A Duplicate the established suppliers' training invoices and change their filenames to the new supplier, preserving the existing verified field labels.

B Representative labeled invoices showing the correct values for the new layout. ✓

C A longer generated explanation appended to each wrong extraction.

D The wrong model outputs relabeled as ground-truth amounts.

Why the correct answer is right

Representative labeled samples can improve an analyzer on new formats. Labels should reflect verified field values from the actual source documents.

Why the other options are wrong

- **A:** The content still lacks the new layout; filenames do not supply representative examples.
- **D:** Incorrect labels reinforce the defect.
- **C:** Explanation length does not supply correct labeled examples.

Sources

Validate document analyzer quality with confidence, grounding, and labeled samples
https://learn.microsoft.com/azure/ai-services/content-understanding/document/analyzer-improvement

Q65 Domain A: Plan and manage an Azure AI solution | Multiple select | hard

A developer can query Search privately, but an agent is denied. The team wants to establish whether the agent has equivalent access. Which TWO observations are directly relevant? Select TWO answers.

A The final answer style requested in the agent's instructions.

B The roles assigned to the developer's interactive account.

C The address and route used from the agent's runtime network. ✓

D The connection principal and its role scope on the Search resource. ✓

E The private endpoint resource ID and subnet assignment, evaluated from the developer's successful network connection to the same service.

Why the correct answer is right

Check the agent's own network path and identity. A developer's successful request does not establish either for a separate runtime.

Why the other options are wrong

- **B:** Those belong to a different caller.
- **E:** Resource configuration and a developer connection do not establish the agent runtime network path.
- **A:** Answer style does not determine service authorization.

Sources

Role-based access control for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry

Set up private networking for Foundry Agent Service
https://learn.microsoft.com/azure/foundry/agents/how-to/virtual-networks

Q66 Domain A: Plan and manage an Azure AI solution | Multiple select | hard

A job can wait for asynchronous results, but its data-processing policy still restricts geography. Which TWO considerations belong in its deployment decision? Select TWO answers.

A Use a resource name containing EU as proof of EU processing.

B Select a supported Batch deployment for the deferred workload. ✓

C Choose Provisioned solely because it supports multiple requests.

D Choose a deployment boundary consistent with the processing policy. ✓

E Use Global Batch because asynchronous completion removes geographic restrictions.

Why the correct answer is right

Completion model and geography remain separate requirements. Batch can fit deferred work, but its boundary must also satisfy the policy.

Why the other options are wrong

- **E:** Completion mode does not waive data policy.
- **A:** Naming does not establish a processing boundary.
- **C:** Concurrency alone does not identify the appropriate batch billing and completion model.

Sources

Deployment types for Microsoft Foundry Models
https://learn.microsoft.com/azure/foundry/foundry-models/concepts/deployment-types

Q67 Domain A: Plan and manage an Azure AI solution | Multiple select | hard

A trace has operation timing and redacted content, but evaluation IDs are missing and access is broader than the logging policy permits. Which TWO changes address these gaps? Select TWO answers.

A Add evaluation-to-request correlation identifiers. ✓

B Create one new trace ID per span to avoid shared identifiers.

C Remove operation timings because content is already redacted.

D Record raw tool credentials to make correlation easier.

E Apply the required trace-storage access and retention controls. ✓

Why the correct answer is right

Correlation connects outcomes to their request traces. Storage access and retention controls apply to traces as they do to operational logs.

Why the other options are wrong

- **D:** Credentials are not correlation identifiers.
- **C:** Timing remains useful diagnostic information.
- **B:** That breaks request relationships.

Sources

Agent tracing overview
https://learn.microsoft.com/azure/foundry/observability/concepts/trace-agent-concept

Q68

Domain A: Plan and manage an Azure AI solution | Multiple select | hard

Two failures occur: session_quota_exceeded with unneeded active sessions, and model-token throttling on sustained traffic. Which TWO diagnoses are correctly scoped? Select TWO answers.

A Review active sessions when addressing the session limit. ✓

B Treat both errors as evidence of invalid Entra credentials.

C Review underlying model throughput for the token throttling. ✓

D Use Search storage expansion as the remedy for both failures.

E Stop sessions as proof that model token throughput has increased.

Why the correct answer is right

Agent-service limits and underlying model quotas coexist. Read the returned code and address the resource implicated by that failure.

Why the other options are wrong

- **D:** Search storage is neither reported constraint.
- **B:** Authentication and quota are different failure classes.
- **E:** Session cleanup does not allocate more model throughput.

Sources

Foundry Agent Service limits, quotas, and regional support
https://learn.microsoft.com/azure/foundry/agents/concepts/limits-quotas-regions

Q69 Domain A: Plan and manage an Azure AI solution | Multiple select | hard

An agent obeys policy but omits requested deliverables and invents one fact absent from retrieved context. Which TWO evaluations directly target the described quality defects? Select TWO answers.

A Tool Input Accuracy as proof of the invented final claim.

B Search latency as the completeness measure.

C Task Adherence as the direct measure of missing deliverables.

D Task Completion (preview) for the missing deliverables. ✓

E Groundedness for contextual support of the generated fact. ✓

Why the correct answer is right

Task completion and groundedness address separate defects: unfinished user outcomes and unsupported generated content.

Why the other options are wrong

- **C:** The scenario states that policy is obeyed; completion is the missing outcome.
- **A:** Argument accuracy does not establish support for generated claims.
- **B:** Timing does not measure answer coverage.

Sources

Agent evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Q70 Domain B: Implement generative AI and agentic solutions | Multiple select | hard

Two tool calls finish concurrently and one result contains a secret in an error message. Which TWO output-handling rules apply? Select TWO answers.

A Use result completion order as the identifier contract.

B Keep the error intact because it came from an internal service.

C Return each result with its originating call_id. ✓

D Associate both outputs with the common function name.

E Sanitize sensitive error content before returning it to the model. ✓

Why the correct answer is right

Preserve the call relationship and minimize sensitive output. These are independent obligations even when execution is concurrent.

Why the other options are wrong

- **D:** The calls require distinct correlation identifiers.
- **A:** Completion order does not replace call_id.
- **B:** Internal results may still contain secrets or untrusted content.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q71 Domain B: Implement generative AI and agentic solutions Multiple select hard

A client moves strict structured output from Chat Completions to Responses. Which TWO implementation details need to be handled correctly? Select TWO answers.

A Remove required nullable fields to reduce the request size.

B Treat JSON parsing success as full compliance with the business schema.

C Place the response format in Responses text.format. ✓

D Handle refusal outcomes separately from successful structured results. ✓

E Keep response_format in the same top-level position as Chat Completions.

Why the correct answer is right

Adapt the API-specific format location and preserve robust handling of non-success outcomes such as refusals.

Why the other options are wrong

- **E:** The destination API uses a different request location.
- **B:** Valid JSON alone does not demonstrate schema adherence.
- **A:** That changes the intended output contract.

Sources

Structured outputs
https://learn.microsoft.com/azure/foundry/openai/how-to/structured-outputs

Q72 Domain B: Implement generative AI and agentic solutions Multiple select hard

A RAG evaluation needs to distinguish search quality from unsupported generation. Labeled query-document judgments and generated answers with context are available. Which TWO assessments fit? Select TWO answers.

A Use answer relevance as proof that all claims are context-supported.

B Evaluate document retrieval against the relevance judgments. ✓

C Use retrieval latency as a replacement for relevance judgments.

D Use token count as the amount of ground-truth support.

E Evaluate groundedness against each answer's supplied context. ✓

Why the correct answer is right

Evaluate each pipeline stage with evidence appropriate to it: labeled retrieval results for search and context-grounded claims for generation.

Why the other options are wrong

- **A:** Relevance and support are different properties.
- **C:** Latency measures operations rather than retrieval quality.
- **D:** Length does not establish evidence support.

Sources

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Q73

Domain B: Implement generative AI and agentic solutions | Multiple select | hard

A purchase agent needs to prevent unsafe actions and inspect receipts returned by the vendor tool. Which TWO intervention points correspond to those requirements? Select TWO answers.

A User-input scanning used as inspection of the later vendor receipt.

B Tool-call inspection before the proposed purchase executes. ✓

C Tool-response inspection before returned vendor content is used. ✓

D Deployment-time validation used as approval for each future purchase.

E Final-output inspection used as prior approval of the purchase.

Why the correct answer is right

Tool-call and tool-response interventions protect different sides of execution. A deployment or final-answer check does not substitute for either boundary.

Why the other options are wrong

- **E:** A final output check cannot approve a completed side effect in advance.
- **A:** The receipt is not present at the original input boundary.
- **D:** Deployment checks do not inspect individual runtime proposals.

Sources

Guardrails and controls overview in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/guardrails/guardrails-overview

Q74 Domain B: Implement generative AI and agentic solutions | Multiple select | medium

A new Foundry Python app uses azure-ai-projects 2.8.0 and needs Responses plus project operations. Which TWO setup choices follow the documented client pattern? Select TWO answers.

A Use a resource key as the AIProjectClient credential.

B Mix classic 1.0.0 project examples into the 2.x request builder unchanged.

C Initialize AIProjectClient with the Foundry project endpoint and Entra credentials. ✓

D Use the Azure OpenAI inference endpoint as the project-management endpoint.

E Obtain its OpenAI client through get_openai_client(). ✓

Why the correct answer is right

Use a coherent 2.x project endpoint and authentication setup, then obtain the supported inference client from AIProjectClient.

Why the other options are wrong

- **D:** The endpoint families expose different operations.
- **A:** The specified client supports Entra authentication.
- **B:** Version-specific examples must match the client contract.

Sources

Azure AI Projects client library for Python - version 2.8.0
https://learn.microsoft.com/python/api/overview/azure/ai-projects-readme?view=azure-python

Microsoft Foundry SDKs and endpoints
https://learn.microsoft.com/azure/foundry/how-to/develop/sdk-overview

Q75

Domain C: Implement computer vision solutions | Multiple select | medium

A site has an image-only download button and a decorative border. Which TWO accessibility choices fit their functions? Select TWO answers.

A Give the button an alternative naming its download action. ✓

B Describe the button's color, outline and position in the header.

C Use the same 'image' label for both elements.

D Omit alt from the decorative border and rely on its filename.

E Use an empty alt attribute for the decorative border. ✓

Why the correct answer is right

Alternative text depends on purpose. Communicate the action for a functional image and use empty alt for a purely decorative image.

Why the other options are wrong

- **B:** Functional alternative text should communicate the download action.
- **D:** An explicit empty alt identifies decorative content.
- **C:** That does not communicate the button function or mark the decoration properly.

Sources

W3C WAI An alt Decision Tree
https://www.w3.org/WAI/tutorials/images/decision-tree/

Q76 Domain C: Implement computer vision solutions | Multiple select | medium

A product offers photograph descriptions and image edits. Which TWO API-design distinctions should it preserve? Select TWO answers.

A All image deployments accept the same mask and quality settings.

B Vision chat receives supported image content alongside text instructions. ✓

C A captioning response is itself the edited image asset.

D Image editing uses the selected image model's supported edit inputs. ✓

E A photograph filename in plain text supplies the image pixels to a remote model.

Why the correct answer is right

Image understanding and image editing have different request and response contracts. Provide actual visual inputs and use model-specific editing capabilities.

Why the other options are wrong

- **E:** A local filename is not image content.
- **C:** Description and image generation are different outputs.
- **A:** Controls vary by model.

Sources

Use vision-enabled chat models
https://learn.microsoft.com/azure/foundry/openai/how-to/gpt-with-vision

Azure OpenAI image generation models
https://learn.microsoft.com/azure/foundry/openai/how-to/dall-e

Q77

Domain C: Implement computer vision solutions | Multiple select | medium

An analyzer uses enableSegment and natural-language contentCategories for customized video analysis. Which TWO planning statements follow the documented behavior? Select TWO answers.

A Segmentation controls the granularity of video analysis. ✓

B Disable segmentation to make sampling include every source frame.

C Segmentation can consume model tokens even when no fields are defined. ✓

D A content category is a command to preserve all original frames.

E Removing all fields makes segmentation token-free.

Why the correct answer is right

Video segmentation and requested categories affect how content is analyzed. Their cost and sampling constraints should be assessed separately.

Why the other options are wrong

- **E:** The documented segmentation process can still consume tokens.
- **D:** Categories are not a lossless frame-sampling control.
- **B:** Disabling segmentation does not remove the documented frame-sampling limitation.

Sources

Azure Content Understanding in Foundry Tools video solutions
https://learn.microsoft.com/azure/ai-services/content-understanding/video/overview

Q78 Domain D: Implement text analysis solutions | Multiple select | medium

A team compares extractive and abstractive summaries for an evidence-sensitive workflow. Which TWO statements are accurate? Select TWO answers.

A Extractive summarization selects sentences from the source. ✓

B Abstractive summarization creates a compact rewording of the content. ✓

C A summarization API converts a source citation into independent verification.

D An abstractive summary's fluent style establishes its factual support.

E An extractive summary necessarily covers every important source point.

Why the correct answer is right

The two methods differ in whether they select or generate wording. Both need evaluation for the workflow's coverage and fidelity requirements.

Why the other options are wrong

- **D:** Fluency does not establish fidelity.
- **E:** Selection can omit important material.
- **C:** Source references support review but do not perform it.

Sources

What is summarization?
https://learn.microsoft.com/azure/ai-services/language-service/summarization/overview

Q79 Domain D: Implement text analysis solutions | Multiple select | medium

A routing pipeline must choose the language-specific queue and locate personal identifiers for redaction. Which TWO Language capabilities directly fit? Select TWO answers.

A Abstractive summarization used as the source language label.

B Language detection for identifying the text's language. ✓

C PII detection for locating personal information. ✓

D General NER used as proof that all personal information was redacted.

E Sentiment analysis used as the language identifier.

Why the correct answer is right

Use task-specific outputs: language detection for routing and PII detection to support redaction. Apply and verify the redaction separately.

Why the other options are wrong

- **E:** Emotional tone is a different output.
- **D:** Entity categories do not establish completion of a PII workflow.
- **A:** A summary is not a language-detection result.

Sources

What is Azure Language in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/language-service/overview

Q80

Domain E: Implement information extraction solutions | Multiple select | medium

An invoice workflow must route uncertain values to reviewers and let them locate the supporting text. Which TWO retained outputs meet those needs? Select TWO answers.

A A model deployment name used as the evidence location.

B The formatted amount stripped of all source metadata.

C A high confidence score treated as an authorization to pay.

D The estimated field confidence for review routing. ✓

E Source page, coordinate or span references for evidence lookup. ✓

Why the correct answer is right

Confidence and grounding serve complementary purposes: deciding what to review and finding the source evidence for that review.

Why the other options are wrong

- **A:** It identifies the model, not the source passage.
- **B:** The value alone cannot locate supporting evidence.
- **C:** Confidence is not business authorization.

Sources

Validate document analyzer quality with confidence, grounding, and labeled samples
https://learn.microsoft.com/azure/ai-services/content-understanding/document/analyzer-improvement

Q81

Domain E: Implement information extraction solutions | Multiple select | hard

Long source documents and persistent embedding throttling block a Search vectorization pipeline. Which TWO adjustments directly address these separate constraints? Select TWO answers.

A Chunk content to fit the embedding model's input limits. ✓

B Replace query vectors with unrelated same-length numerical arrays.

C Increase chat response length to enlarge embedding input capacity.

D Align ingestion demand with available embedding-model quota. ✓

E Retry the unchanged workload indefinitely to create more quota.

Why the correct answer is right

Input size and throughput quota are different embedding constraints. Chunking addresses document size; capacity or scheduling changes address sustained quota exhaustion.

Why the other options are wrong

- **C:** Chat generation limits do not change embedding limits.
- **E:** Retries do not allocate persistent capacity.
- **B:** Arbitrary vectors do not preserve semantic retrieval.

Sources

Integrated vector embedding in Azure AI Search
https://learn.microsoft.com/azure/search/vector-search-integrated-vectorization

Q82 Domain E: Implement information extraction solutions | Multiple select | hard

A Search solution must index skill-generated text and then combine exact-term and semantic retrieval. Which TWO configurations belong at the appropriate stages? Select TWO answers.

A Use outputFieldMappings for the enriched text's index field. ✓

B Use hybrid search to combine keyword and vector result lists. ✓

C Use a query filter to generate document embeddings during ingestion.

D Use semantic reranking as the source-to-index field mapper.

E Use fieldMappings to read an enrichment node as an original source field.

Why the correct answer is right

Separate ingestion and retrieval responsibilities. Output mappings populate enriched fields, while hybrid retrieval combines complementary ranking signals.

Why the other options are wrong

- **E:** The node was generated by a skill.
- **D:** Reranking is a query stage.
- **C:** Filters do not create stored vectors.

Sources

Hybrid search using vectors and full-text search in Azure AI Search
https://learn.microsoft.com/azure/search/hybrid-search-overview

Skillset concepts in Azure AI Search
https://learn.microsoft.com/azure/search/cognitive-search-working-with-skillsets

Q83 Domain A: Plan and manage an Azure AI solution Sequence medium

A support run has already been evaluated and its score is low. Arrange the prescribed investigation: start from the evaluation record, locate its trace, inspect the slow child operation, then record the diagnosis.

Arrange these steps in order:

A. Locate the corresponding request trace.

B. Inspect the timed child spans to identify the slow operation.

C. Read the evaluation run identifier and linked request identifier.

D. Record the operation-level diagnosis with its correlation references.

Correct order:

1. C Read the evaluation run identifier and linked request identifier.

2. A Locate the corresponding request trace.

3. B Inspect the timed child spans to identify the slow operation.

4. D Record the operation-level diagnosis with its correlation references.

Why the correct answer is right

The identifiers connect the evaluation to the request; nested timing then identifies the operation supporting the diagnosis.

Sources

Agent tracing overview
https://learn.microsoft.com/azure/foundry/observability/concepts/trace-agent-concept

Q84 Domain A: Plan and manage an Azure AI solution Sequence medium

A private endpoint exists. Troubleshooting policy requires confirming DNS, then network reachability, then checking authorization after an actual denial. Arrange this investigation.

Arrange these steps in order:

A. Inspect the denied request to identify its runtime connection principal.

B. Check that principal's required role assignments on the target resource.

C. Confirm connectivity to the resulting private endpoint address.

D. Resolve the service hostname from the agent runtime network.

Correct order:

1. D Resolve the service hostname from the agent runtime network.

2. C Confirm connectivity to the resulting private endpoint address.

3. A Inspect the denied request to identify its runtime connection principal.

4. B Check that principal's required role assignments on the target resource.

Why the correct answer is right

This prescribed procedure separates DNS and network access from identity authorization. Developer access is not substituted for runtime evidence.

Sources

Role-based access control for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry

Set up private networking for Foundry Agent Service
https://learn.microsoft.com/azure/foundry/agents/how-to/virtual-networks

Q85 Domain A: Plan and manage an Azure AI solution Sequence medium

A deployment decision follows the company checklist in this order: geography, completion mode, billing, then supported model availability. Arrange the checks for a deferred document-classification job.

Arrange these steps in order:

A. Check that the required model and capacity are available for that choice.

B. Confirm that asynchronous batch completion fits the job.

C. Select the compatible batch charging and deployment choice.

D. Identify the permitted inference processing boundary.

Correct order:

1. D Identify the permitted inference processing boundary.

2. B Confirm that asynchronous batch completion fits the job.

3. C Select the compatible batch charging and deployment choice.

4. A Check that the required model and capacity are available for that choice.

Why the correct answer is right

The ordering is specified by the scenario checklist. Deployment type decisions combine geography, completion and charging, followed by model availability checks.

Sources

Deployment types for Microsoft Foundry Models
https://learn.microsoft.com/azure/foundry/foundry-models/concepts/deployment-types

Q86 Domain B: Implement generative AI and agentic solutions Sequence medium

A read-only function call has been returned by Responses. Arrange this continuation path: execute validated input, sanitize the result, return it with its identifier, then continue generation.

Arrange these steps in order:

A. Continue the model response using the correlated function output.

B. Submit function_call_output with the originating call_id.

C. Validate the arguments and execute the authorized read-only function.

D. Remove sensitive fields from the returned tool data.

Correct order:

1. C Validate the arguments and execute the authorized read-only function.

2. D Remove sensitive fields from the returned tool data.

3. B Submit function_call_output with the originating call_id.

4. A Continue the model response using the correlated function output.

Why the correct answer is right

The host obtains an actual result, sanitizes it before model exposure, preserves its call association and then continues the response.

Sources

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q87 Domain B: Implement generative AI and agentic solutions Sequence medium

A receipt assistant follows this prescribed request flow: receive an uploaded receipt, retrieve its evidence, draft a proposed refund, then obtain approval before execution. Arrange the stages.

Arrange these steps in order:

A. Use file search to retrieve relevant receipt evidence.

B. Prepare and validate a refund proposal using that evidence.

C. Obtain approval before executing the proposed refund.

D. Receive the user-provided receipt for the interaction.

Correct order:

1. D Receive the user-provided receipt for the interaction.

2. A Use file search to retrieve relevant receipt evidence.

3. B Prepare and validate a refund proposal using that evidence.

4. C Obtain approval before executing the proposed refund.

Why the correct answer is right

File search supplies evidence for the proposal; evidence is not authorization, so the side effect still requires confirmation.

Sources

Memory in Microsoft Foundry Agent Service (preview)
https://learn.microsoft.com/azure/foundry/agents/concepts/what-is-memory

Use function calling with Microsoft Foundry agents
https://learn.microsoft.com/azure/foundry/agents/how-to/tools/function-calling

Q88 Domain C: Implement computer vision solutions Sequence medium

An image deployment already supports the requested edit. Arrange this preparation workflow: load the source, build a matching mask, submit the edit, then inspect its result.

Arrange these steps in order:

A. Submit the source, mask and edit prompt to the supported deployment.

B. Prepare the edit mask with the required matching dimensions.

C. Load the source image and determine its dimensions.

D. Inspect the returned image against the requested edit and preserved region.

Correct order:

1. C Load the source image and determine its dimensions.

2. B Prepare the edit mask with the required matching dimensions.

3. A Submit the source, mask and edit prompt to the supported deployment.

4. D Inspect the returned image against the requested edit and preserved region.

Why the correct answer is right

Mask preparation depends on the source dimensions. The result can be inspected after the edit request completes.

Sources

Azure OpenAI image generation models
https://learn.microsoft.com/azure/foundry/openai/how-to/dall-e

Q89 Domain D: Implement text analysis solutions | Sequence | medium

A successful speech translation result is available for German. Arrange the client's prescribed playback path: choose target text, submit it to synthesis, obtain audio, then play it.

Arrange these steps in order:

A. Play the produced audio through the client output device.

B. Select the de entry from the translation result.

C. Submit that German text to the configured SpeechSynthesizer.

D. Receive the successful synthesis audio result.

Correct order:

1. B Select the de entry from the translation result.

2. C Submit that German text to the configured SpeechSynthesizer.

3. D Receive the successful synthesis audio result.

4. A Play the produced audio through the client output device.

Why the correct answer is right

The target text is the synthesis input; playback depends on the produced audio. Translation success alone does not mean audio was synthesized.

Sources

How to recognize and translate speech
https://learn.microsoft.com/azure/ai-services/speech-service/how-to-translate-speech

Q90 Domain D: Implement text analysis solutions | Sequence | medium

Company policy requires language routing, then redaction, then translation of the redacted text. Arrange this text-processing workflow.

Arrange these steps in order:

A. Apply the required redactions using those detected spans.

B. Detect personal-information spans in the routed text.

C. Translate the approved redacted text into the target language.

D. Detect the submitted text's language for routing.

Correct order:

1. D Detect the submitted text's language for routing.

2. B Detect personal-information spans in the routed text.

3. A Apply the required redactions using those detected spans.

4. C Translate the approved redacted text into the target language.

Why the correct answer is right

The ordering follows the stated company policy. Language detection, PII detection, redaction and translation are separate operations.

Sources

What is Azure Language in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/language-service/overview

What is Azure Translator in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/translator/overview

Q91 Domain E: Implement information extraction solutions Sequence medium

The OCR skill and index schema are configured. Arrange the document data path from normalized images through searchable enriched text.

Arrange these steps in order:

A. Pass those image inputs through the OCR skill.

B. Obtain the OCR text in the enriched document tree.

C. Map the enriched OCR text into the searchable index field.

D. Produce normalized images through the blob indexer imageAction setting.

Correct order:

1. D Produce normalized images through the blob indexer imageAction setting.

2. A Pass those image inputs through the OCR skill.

3. B Obtain the OCR text in the enriched document tree.

4. C Map the enriched OCR text into the searchable index field.

Why the correct answer is right

The OCR skill needs normalized images. Its enriched text output must then be mapped into the target index.

Sources

OCR cognitive skill
https://learn.microsoft.com/azure/search/cognitive-search-skill-ocr

Skillset concepts in Azure AI Search
https://learn.microsoft.com/azure/search/cognitive-search-working-with-skillsets

Q92 Domain E: Implement information extraction solutions Sequence medium

An invoice field has low confidence. Company policy requires evidence inspection and confirmation before posting. Arrange that review workflow.

Arrange these steps in order:

A. Open the field's referenced source page and region.

B. Route the uncertain field to a human review queue.

C. Confirm or correct the value using the source evidence.

D. Post the approved value to the downstream system.

Correct order:

1. B Route the uncertain field to a human review queue.

2. A Open the field's referenced source page and region.

3. C Confirm or correct the value using the source evidence.

4. D Post the approved value to the downstream system.

Why the correct answer is right

Confidence triggers the review path; grounding locates the evidence. Posting follows the required human confirmation.

Sources

Validate document analyzer quality with confidence, grounding, and labeled samples
https://learn.microsoft.com/azure/ai-services/content-understanding/document/analyzer-improvement

Q93 Domain A: Plan and manage an Azure AI solution | Matching | medium

Match each observed agent limit to the next listed response. Retained session state must be preserved; regional moves are prohibited. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. Unneeded active sessions cause session_quota_exceeded
Match: _____

B. Transient regional_session_quota_exceeded
Match: _____

C. Regional session failure persists after suitable retries
Match: _____

Choice bank

- Increase Search storage
- Use exponential backoff with jitter
- Stop the unneeded active sessions
- Open an Azure support request

Correct matches:

A. Unneeded active sessions cause session_quota_exceeded	**→ Stop the unneeded active sessions**
B. Transient regional_session_quota_exceeded	**→ Use exponential backoff with jitter**
C. Regional session failure persists after suitable retries	**→ Open an Azure support request**

Why the correct answer is right

The remedies follow the returned limit and scenario constraints. Search storage is unrelated to these session-capacity failures.

Sources

Foundry Agent Service limits, quotas, and regional support
https://learn.microsoft.com/azure/foundry/agents/concepts/limits-quotas-regions

Q94 Domain A: Plan and manage an Azure AI solution | Matching | medium

Match each design requirement to the configuration area that controls it. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. Agent state stored in customer-managed resources
Match: _____

B. Private service hostname resolution from the agent network
Match: _____

C. Scoped permissions enforced for service calls
Match: _____

Choice bank

- Model output temperature
- Private DNS and network configuration
- Entra authentication with RBAC
- Standard agent setup

Correct matches:

A. Agent state stored in customer-managed resources	**→ Standard agent setup**
B. Private service hostname resolution from the agent network	**→ Private DNS and network configuration**
C. Scoped permissions enforced for service calls	**→ Entra authentication with RBAC**

Why the correct answer is right

Storage setup, network resolution and identity permissions solve separate requirements. Output temperature does not configure them.

Sources

Foundry Agent Service limits, quotas, and regional support
https://learn.microsoft.com/azure/foundry/agents/concepts/limits-quotas-regions

Role-based access control for Microsoft Foundry
https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry

Set up private networking for Foundry Agent Service
https://learn.microsoft.com/azure/foundry/agents/how-to/virtual-networks

Q95 Domain B: Implement generative AI and agentic solutions Matching medium

Match each evaluation question to the most directly relevant measure. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. Are the answer's claims supported by its supplied passages?
Match: _____

B. Does the answer address the user's question?
Match: _____

C. Did the agent deliver the requested overall outcome?
Match: _____

Choice bank

- Task Completion (preview)
- Search latency
- Relevance
- Groundedness

Correct matches:

A. Are the answer's claims supported by its supplied passages?	**→ Groundedness**
B. Does the answer address the user's question?	**→ Relevance**
C. Did the agent deliver the requested overall outcome?	**→ Task Completion (preview)**

Why the correct answer is right

Contextual support, response relevance and task completion are different quality properties. Search latency is operational timing.

Sources

Retrieval-Augmented Generation (RAG) evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/rag-evaluators

Agent evaluators
https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators

Q96 Domain B: Implement generative AI and agentic solutions Matching medium

Match each content boundary to its agent guardrail intervention point. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. A proposed purchase action before execution
Match: _____

B. Data returned by the vendor service before agent use
Match: _____

C. The completed answer before user delivery
Match: _____

Choice bank

- Tool response
- Tool call
- Final output
- User input

Correct matches:

A. A proposed purchase action before execution	**→ Tool call**
B. Data returned by the vendor service before agent use	**→ Tool response**
C. The completed answer before user delivery	**→ Final output**

Why the correct answer is right

The controls align with the moment data or action crosses a boundary. User-input inspection happens before these later artifacts exist.

Sources

Guardrails and controls overview in Microsoft Foundry
https://learn.microsoft.com/azure/foundry/guardrails/guardrails-overview

Q97 Domain C: Implement computer vision solutions | Matching | medium

Match each visual-content requirement to the relevant capability. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. Assess visual harm categories in an uploaded image
Match: _____

B. Describe objects in a supplied photograph
Match: _____

C. Assess malicious instructions in extracted poster text
Match: _____

Choice bank

- Vision-enabled chat
- Custom Speech recognition
- Text Prompt Shields
- Image Content Safety analysis

Correct matches:

A. Assess visual harm categories in an uploaded image	**→ Image Content Safety analysis**
B. Describe objects in a supplied photograph	**→ Vision-enabled chat**
C. Assess malicious instructions in extracted poster text	**→ Text Prompt Shields**

Why the correct answer is right

Visual description, image harm classification and text instruction security are distinct tasks with different input contracts.

Sources

QuickStart: Analyze image content
https://learn.microsoft.com/azure/ai-services/content-safety/quickstart-image

Azure AI Content Safety Frequently Asked Questions
https://learn.microsoft.com/en-us/azure/ai-services/content-safety/faq

Use vision-enabled chat models
https://learn.microsoft.com/azure/foundry/openai/how-to/gpt-with-vision

Q98 Domain D: Implement text analysis solutions | Matching | medium

Match each requested output to the relevant text-processing approach. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. A summary made from selected source sentences
Match: _____

B. A concise reworded synthesis of document ideas
Match: _____

C. The document's content in another language
Match: _____

Choice bank

- Extractive summarization
- Document translation
- Abstractive summarization
- Language detection

Correct matches:

A. A summary made from selected source sentences	**→ Extractive summarization**
B. A concise reworded synthesis of document ideas	**→ Abstractive summarization**
C. The document's content in another language	**→ Document translation**

Why the correct answer is right

Extraction preserves selected sentences; abstraction generates new wording; translation changes language. Detecting a language does not produce those outputs.

Sources

What is summarization?
https://learn.microsoft.com/azure/ai-services/language-service/summarization/overview

What is Azure Translator in Foundry Tools?
https://learn.microsoft.com/azure/ai-services/translator/overview

Q99 Domain D: Implement text analysis solutions | Matching | medium

Match each Italian-to-German speech workflow setting or value to its purpose. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. it-IT on speech_recognition_language
Match: _____

B. de added as the translation target
Match: _____

C. The successful de translation string sent to SpeechSynthesizer
Match: _____

Choice bank

- Requested output language
- Text input for synthesis
- Spoken source language
- Recognition endpoint credential

Correct matches:

A. it-IT on speech_recognition_language	**→ Spoken source language**
B. de added as the translation target	**→ Requested output language**
C. The successful de translation string sent to SpeechSynthesizer	**→ Text input for synthesis**

Why the correct answer is right

Language identifiers and the translated text have different roles. None of these values is a credential.

Sources

How to recognize and translate speech
https://learn.microsoft.com/azure/ai-services/speech-service/how-to-translate-speech

Q100 Domain E: Implement information extraction solutions | Matching | medium

Match each extraction task to its most direct mechanism. Use each selected choice once; one choice is unused.

Match each row to a choice.

A. Rename an original source field during ingestion
Match: _____

B. Route an uncertain extracted value for review
Match: _____

C. Track an analysis request until results are ready
Match: _____

Choice bank

- Long-running operation poller
- fieldMappings
- Estimated field confidence
- Semantic reranking

Correct matches:

A. Rename an original source field during ingestion	**→ fieldMappings**
B. Route an uncertain extracted value for review	**→ Estimated field confidence**
C. Track an analysis request until results are ready	**→ Long-running operation poller**

Why the correct answer is right

Source mapping, uncertainty routing and asynchronous operation tracking solve different problems. Reranking applies to retrieved candidates.

Sources

Skillset concepts in Azure AI Search
https://learn.microsoft.com/azure/search/cognitive-search-working-with-skillsets

Validate document analyzer quality with confidence, grounding, and labeled samples
https://learn.microsoft.com/azure/ai-services/content-understanding/document/analyzer-improvement

Azure AI Content Understanding client library for Python - version 1.1.0
https://learn.microsoft.com/python/api/overview/azure/ai-contentunderstanding-readme?view=azure-python

Quick Answer Key

#	Type	Correct Answer
Q1	Single choice	**C**
Q2	Single choice	**D**
Q3	Single choice	**D**
Q4	Single choice	**C**
Q5	Single choice	**A**
Q6	Single choice	**B**
Q7	Single choice	**D**
Q8	Single choice	**B**
Q9	Single choice	**B**
Q10	Single choice	**B**
Q11	Single choice	**B**
Q12	Single choice	**A**
Q13	Single choice	**A**
Q14	Single choice	**C**
Q15	Single choice	**A**
Q16	Single choice	**B**
Q17	Single choice	**D**
Q18	Single choice	**B**
Q19	Case study	**C**
Q20	Case study	**D**
Q21	Case study	**B**
Q22	Case study	**D**
Q23	Case study	**B**
Q24	Case study	**C**
Q25	Case study	**A**
Q26	Case study	**B**
Q27	Single choice	**D**
Q28	Configuration	**A**
Q29	Single choice	**A**
Q30	Single choice	**D**
Q31	Single choice	**D**
Q32	Single choice	**C**
Q33	Single choice	**D**
Q34	Single choice	**A**
Q35	Single choice	**B**

#	Type	Correct Answer
Q36	Single choice	**C**
Q37	Single choice	**D**
Q38	Single choice	**C**
Q39	Single choice	**D**
Q40	Single choice	**C**
Q41	Single choice	**B**
Q42	Single choice	**C**
Q43	Single choice	**C**
Q44	Single choice	**C**
Q45	Single choice	**D**
Q46	Single choice	**A**
Q47	Configuration	**C**
Q48	Single choice	**A**
Q49	Single choice	**D**
Q50	Single choice	**A**
Q51	Single choice	**C**
Q52	Single choice	**A**
Q53	Single choice	**B**
Q54	Single choice	**A**
Q55	Single choice	**B**
Q56	Configuration	**A**
Q57	Single choice	**C**
Q58	Single choice	**D**
Q59	Single choice	**A**
Q60	Single choice	**B**
Q61	Single choice	**D**
Q62	Single choice	**A**
Q63	Single choice	**C**
Q64	Single choice	**B**
Q65	Multiple select	**C, D**
Q66	Multiple select	**B, D**
Q67	Multiple select	**A, E**
Q68	Multiple select	**A, C**
Q69	Multiple select	**D, E**
Q70	Multiple select	**C, E**
Q71	Multiple select	**C, D**

#	Type	Correct Answer
Q72	Multiple select	**B, E**
Q73	Multiple select	**B, C**
Q74	Multiple select	**C, E**
Q75	Multiple select	**A, E**
Q76	Multiple select	**B, D**
Q77	Multiple select	**A, C**
Q78	Multiple select	**A, B**
Q79	Multiple select	**B, C**
Q80	Multiple select	**D, E**
Q81	Multiple select	**A, D**
Q82	Multiple select	**A, B**
Q83	Sequence	**1. C → 2. A → 3. B → 4. D**
Q84	Sequence	**1. D → 2. C → 3. A → 4. B**
Q85	Sequence	**1. D → 2. B → 3. C → 4. A**
Q86	Sequence	**1. C → 2. D → 3. B → 4. A**
Q87	Sequence	**1. D → 2. A → 3. B → 4. C**
Q88	Sequence	**1. C → 2. B → 3. A → 4. D**
Q89	Sequence	**1. B → 2. C → 3. D → 4. A**
Q90	Sequence	**1. D → 2. B → 3. A → 4. C**
Q91	Sequence	**1. D → 2. A → 3. B → 4. C**
Q92	Sequence	**1. B → 2. A → 3. C → 4. D**
Q93	Matching	**A → Stop the unneeded active sessions | B → Use exponential backoff with jitter | C → Open an Azure support request**
Q94	Matching	**A → Standard agent setup | B → Private DNS and network configuration | C → Entra authentication with RBAC**
Q95	Matching	**A → Groundedness | B → Relevance | C → Task Completion (preview)**
Q96	Matching	**A → Tool call | B → Tool response | C → Final output**
Q97	Matching	**A → Image Content Safety analysis | B → Vision-enabled chat | C → Text Prompt Shields**
Q98	Matching	**A → Extractive summarization | B → Abstractive summarization | C → Document translation**
Q99	Matching	**A → Spoken source language | B → Requested output language | C → Text input for synthesis**
Q100	Matching	**A → fieldMappings | B → Estimated field confidence | C → Long-running operation poller**